

Set zone

Set your working Google Cloud zone by running the following command, substituting the local zone as `us-central1-a`:

```
gcloud config set compute/zone us-central1-a
```

Get sample code for this lab

Get the sample code for creating and running containers and deployments:

```
gsutil -m cp -r gs://spls/gsp053/orchestrate-with-kubernetes .  
cd orchestrate-with-kubernetes/kubernetes
```

Create a cluster with five `n1-standard-1` nodes (this will take a few minutes to complete):

```
gcloud container clusters create bootcamp --num-nodes 5 --scopes  
"https://www.googleapis.com/auth/projecthosting,storage-rw"
```

Learn about the deployment object

Let's get started with Deployments. First let's take a look at the Deployment object. The `explain` command in `kubectl` can tell us about the Deployment object.

```
kubectl explain deployment
```

We can also see all of the fields using the `--recursive` option.

```
kubectl explain deployment --recursive
```

You can use the `explain` command as you go through the lab to help you understand the structure of a Deployment object and understand what the individual fields do.

```
kubectl explain deployment.metadata.name
```

Create a deployment

Update the `deployments/auth.yaml` configuration file:

```
vi deployments/auth.yaml
```

Start the editor:

```
i
```

Change the `image` in the `containers` section of the Deployment to the following:

```
...
containers:
- name: auth
  image: "kelseyhightower/auth:1.0.0"
...
```

Save the `auth.yaml` file: press `<Esc>` then type:

```
:wq
```

Press `<Enter>`. Now let's create a simple deployment. Examine the deployment configuration file:

```
cat deployments/auth.yaml
```

(Output)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: auth
spec:
  replicas: 1
  selector:
    matchLabels:
      app: auth
  template:
    metadata:
      labels:
        app: auth
        track: stable
    spec:
      containers:
      - name: auth
        image: "kelseyhightower/auth:1.0.0"
        ports:
        - name: http
          containerPort: 80
        - name: health
          containerPort: 81
      ...
```

Notice how the Deployment is creating one replica and it's using version 1.0.0 of the `auth` container.

When you run the `kubectl create` command to create the auth deployment, it will make one pod that conforms to the data in the Deployment manifest. This means we can scale the number of Pods by changing the number specified in the `replicas` field.

Go ahead and create your deployment object using `kubectl create`:

```
kubectl create -f deployments/auth.yaml
```

Once you have created the Deployment, you can verify that it was created.

```
kubectl get deployments
```

Once the deployment is created, Kubernetes will create a ReplicaSet for the Deployment. We can verify that a ReplicaSet was created for our Deployment:

```
kubectl get replicaset
```

We should see a ReplicaSet with a name like `auth-xxxxxxx`

Finally, we can view the Pods that were created as part of our Deployment. The single Pod is created by the Kubernetes when the ReplicaSet is created.

```
kubectl get pods
```

It's time to create a service for our auth deployment. You've already seen service manifest files, so we won't go into the details here. Use the `kubectl create` command to create the auth service.

```
kubectl create -f services/auth.yaml
```

Now, do the same thing to create and expose the hello Deployment.

```
kubectl create -f deployments/hello.yaml
```

```
kubectl create -f services/hello.yaml
```

And one more time to create and expose the frontend Deployment.

```
kubectl create secret generic tls-certs --from-file tls/
```

```
kubectl create configmap nginx-frontend-conf --from-file=nginx/frontend.conf
```

```
kubectl create -f deployments/frontend.yaml
```

```
kubectl create -f services/frontend.yaml
```

Note: You created a ConfigMap for the frontend.

Interact with the frontend by grabbing its external IP and then curling to it.

```
kubectl get services frontend
```

It may take a few seconds before the External-IP field is populated for your service. This is normal. Just re-run the above command every few seconds until the field is populated.

```
curl -ks https://<EXTERNAL-IP>
```

And you get the hello response back.

You can also use the output templating feature of `kubectl` to use curl as a one-liner:

```
curl -ks https://`kubectl get svc frontend -o=jsonpath='{.status.loadBalancer.ingress[0].ip}'`
```

Test Completed Task

Click **Check my progress** below to check your lab progress. If you successfully created Kubernetes cluster and Auth, Hello and Frontend deployments, you'll see an assessment score.

Create a Kubernetes cluster and deployments (Auth, Hello, and Frontend)

[Check my progress](#)

Scale a Deployment

Now that we have a Deployment created, we can scale it. Do this by updating the `spec.replicas` field. You can look at an explanation of this field using the `kubectl explain` command again.

```
kubectl explain deployment.spec.replicas
```

The `replicas` field can be most easily updated using the `kubectl scale` command:

```
kubectl scale deployment hello --replicas=5
```

Note: It may take a minute or so for all the new pods to start up.

After the Deployment is updated, Kubernetes will automatically update the associated ReplicaSet and start new Pods to make the total number of Pods equal 5.

Verify that there are now 5 `hello` Pods running:

```
kubectl get pods | grep hello- | wc -l
```

Now scale back the application:

```
kubectl scale deployment hello --replicas=3
```

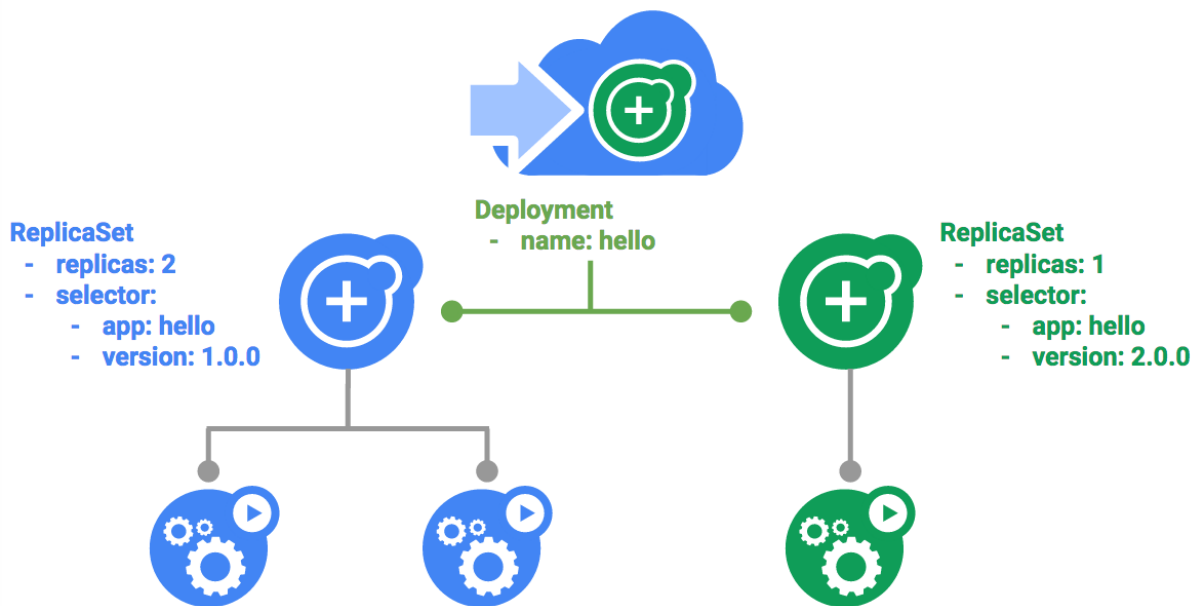
Again, verify that you have the correct number of Pods.

```
kubectl get pods | grep hello- | wc -l
```

You learned about Kubernetes deployments and how to manage & scale a group of Pods.

Rolling update

Deployments support updating images to a new version through a rolling update mechanism. When a Deployment is updated with a new version, it creates a new ReplicaSet and slowly increases the number of replicas in the new ReplicaSet as it decreases the replicas in the old ReplicaSet.



Trigger a rolling update

To update your Deployment, run the following command:

```
kubectl edit deployment hello
```

Change the `image` in the `containers` section of the Deployment to the following:

```
...  
containers:  
  image: kelseyhightower/hello:2.0.0  
...
```

Save and exit.

Once you save out of the editor, the updated Deployment will be saved to your cluster and Kubernetes will begin a rolling update.

See the new ReplicaSet that Kubernetes creates.:

```
kubectl get replicaset
```

You can also see a new entry in the rollout history:

```
kubectl rollout history deployment/hello
```

Pause a rolling update

If you detect problems with a running rollout, pause it to stop the update. Give that a try now:

```
kubectl rollout pause deployment/hello
```

Verify the current state of the rollout:

```
kubectl rollout status deployment/hello
```

You can also verify this on the Pods directly:

```
kubectl get pods -o jsonpath --template='{range .items[*]}.{.metadata.name} {"\t"} {"\t"} {.spec.containers[0].image} {"\n"} {end }'
```

Resume a rolling update

The rollout is paused which means that some pods are at the new version and some pods are at the older version. We can continue the rollout using the `resume` command.

```
kubectl rollout resume deployment/hello
```

When the rollout is complete, you should see the following when running the `status` command.

```
kubectl rollout status deployment/hello
```

(Output)

```
deployment "hello" successfully rolled out
```

Rollback an update

Assume that a bug was detected in your new version. Since the new version is presumed to have problems, any users connected to the new Pods will experience those issues.

You will want to roll back to the previous version so you can investigate and then release a version that is fixed properly.

Use the `rollout` command to roll back to the previous version:

```
kubectl rollout undo deployment/hello
```

Verify the roll back in the history:

```
kubectl rollout history deployment/hello
```

Finally, verify that all the Pods have rolled back to their previous versions:

```
kubectl get pods -o jsonpath --template='{range .items[*]}.{.metadata.name}{"\t"}{"\t"}{.spec.containers[0].image}{"\n"}{end }'
```

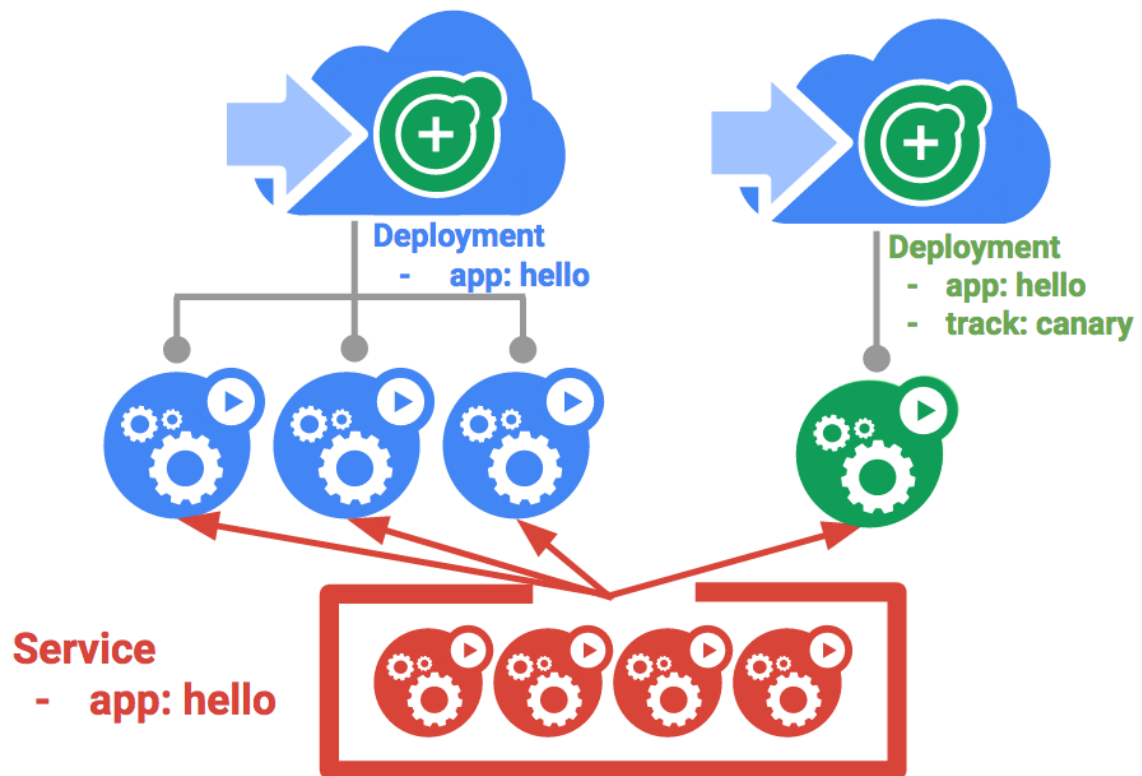
Great! You learned about rolling updates for Kubernetes deployments and how to update applications without downtime.

Canary deployments

When you want to test a new deployment in production with a subset of your users, use a canary deployment. Canary deployments allow you to release a change to a small subset of your users to mitigate risk associated with new releases.

Create a canary deployment

A canary deployment consists of a separate deployment with your new version and a service that targets both your normal, stable deployment as well as your canary deployment.



First, create a new canary deployment for the new version:

```
cat deployments/hello-canary.yaml
```

(Output)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-canary
spec:
  replicas: 1
  selector:
    matchLabels:
      app: hello
  template:
    metadata:
      labels:
        app: hello
        track: canary
        # Use ver 2.0.0 so it matches version on service selector
        version: 2.0.0
    spec:
      containers:
        - name: hello
          image: kelseyhightower/hello:2.0.0
          ports:
            - name: http
              containerPort: 80
            - name: health
              containerPort: 81
...

```

Now create the canary deployment:


```
kubectl create -f deployments/hello-canary.yaml
```

After the canary deployment is created, you should have two deployments, `hello` and `hello-canary`. Verify it with this `kubectl` command:

```
kubectl get deployments
```

On the `hello` service, the selector uses the `app:hello` selector which will match pods in **both** the prod deployment and canary deployment. However, because the canary deployment has a fewer number of pods, it will be visible to fewer users.

Verify the canary deployment

You can verify the `hello` version being served by the request:

```
curl -ks https://`kubectl get svc frontend -o=jsonpath='{.status.loadBalancer.ingress[0].ip}'`/version
```

Run this several times and you should see that some of the requests are served by `hello 1.0.0` and a small subset ($1/4 = 25\%$) are served by `2.0.0`.

Test Completed Task

Click **Check my progress** below to check your lab progress. If you successfully created Canary deployment, you'll see an assessment score.

Canary Deployment

Check my progress

Canary deployments in production - session affinity

In this lab, each request sent to the Nginx service had a chance to be served by the canary deployment. But what if you wanted to ensure that a user didn't get served by the Canary deployment? A use case could be that the UI for an application changed, and you don't want to confuse the user. In a case like this, you want the user to "stick" to one deployment or the other.

You can do this by creating a service with session affinity. This way the same user will always be served from the same version. In the example below the service is the same as before, but a new `sessionAffinity` field has been added, and set to `ClientIP`. All clients with the same IP address will have their requests sent to the same version of the `hello` application.

```
kind: Service
apiVersion: v1
```


```
metadata:
  name: "hello"
spec:
  sessionAffinity: ClientIP
  selector:
    app: "hello"
  ports:
    - protocol: "TCP"
      port: 80
      targetPort: 80
```

Due to it being difficult to set up an environment to test this, you don't need to here, but you may want to use `sessionAffinity` for canary deployments in production.

Blue-green deployments

Rolling updates are ideal because they allow you to deploy an application slowly with minimal overhead, minimal performance impact, and minimal downtime. There are instances where it is beneficial to modify the load balancers to point to that new version only after it has been fully deployed. In this case, blue-green deployments are the way to go.

Kubernetes achieves this by creating two separate deployments; one for the old "blue" version and one for the new "green" version. Use your existing `hello` deployment for the "blue" version. The deployments will be accessed via a Service which will act as the router. Once the new "green" version is up and running, you'll switch over to using that version by updating the Service.



A major downside of blue-green deployments is that you will need to have at least 2x the resources in your cluster necessary to host your application. Make sure you have enough resources in your cluster before deploying both versions of the application at once.

The service

Use the existing hello service, but update it so that it has a selector `app:hello, version: 1.0.0`. The selector will match the existing "blue" deployment. But it will not match the "green" deployment because it will use a different version.

First update the service:

```
kubectl apply -f services/hello-blue.yaml
```

NOTE: Ignore the warning that says resource `service/hello` is missing as this is patched automatically.

Updating using Blue-Green Deployment

In order to support a blue-green deployment style, we will create a new "green" deployment for our new version. The green deployment updates the version label and the image path.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-green
spec:
  replicas: 3
  selector:
    matchLabels:
      app: hello
  template:
    metadata:
      labels:
        app: hello
        track: stable
        version: 2.0.0
    spec:
      containers:
        - name: hello
          image: kelseyhightower/hello:2.0.0
          ports:
            - name: http
              containerPort: 80
            - name: health
              containerPort: 81
          resources:
            limits:
              cpu: 0.2
              memory: 10Mi
          livenessProbe:
            httpGet:
              path: /healthz
              port: 81
              scheme: HTTP
            initialDelaySeconds: 5
            periodSeconds: 15
            timeoutSeconds: 5
          readinessProbe:
            httpGet:
              path: /readiness
              port: 81
              scheme: HTTP
            initialDelaySeconds: 5
            timeoutSeconds: 1
```

Create the green deployment:

```
kubectl create -f deployments/hello-green.yaml
```

Once you have a green deployment and it has started up properly, verify that the current version of 1.0.0 is still being used:

```
curl -ks https://`kubectl get svc frontend -o=jsonpath='{.status.loadBalancer.ingress[0].ip}'`/version
```

Now, update the service to point to the new version:

```
kubectl apply -f services/hello-green.yaml
```

With the service is updated, the "green" deployment will be used immediately. You can now verify that the new version is always being used.

```
curl -ks https://`kubectl get svc frontend -o=jsonpath='{.status.loadBalancer.ingress[0].ip}'`/version
```

Blue-Green Rollback

If necessary, you can roll back to the old version in the same way. While the "blue" deployment is still running, just update the service back to the old version.

```
kubectl apply -f services/hello-blue.yaml
```

Once you have updated the service, your rollback will have been successful. Again, verify that the right version is now being used:

```
curl -ks https://`kubectl get svc frontend -o=jsonpath='{.status.loadBalancer.ingress[0].ip}'`/version
```

You did it! You learned about blue-green deployments and how to deploy updates to applications that need to switch versions all at once.